

TensorFlow GPU @AWS

Sung Kim <hunkim+ml@gmail.com>
<http://hunkim.github.io/ml/>

Deep Network

- Takes a long time for training
 - Many forward/backward propagation and weight updates
 - Many metrics multiplications
- Very quick for testing and use in practice
 - One simple forward propagation

GPU

- A graphics processing unit (GPU), also occasionally called visual processing unit (VPU), is a specialized electronic circuit designed to rapidly manipulate and alter memory to accelerate the creation of images in a frame buffer intended for output to a display.



GPU version

```
# Ubuntu/Linux 64-bit, GPU enabled, Python 2.7
# Requires CUDA toolkit 7.5 and CuDNN v4. For other versions, see "Install from
sources" below.
$ export TF_BINARY_URL=https://storage.googleapis.com/tensorflow/linux/gpu/tensorflow-0.9.0rc0-cp27-none-linux\_x86\_64.whl

# Python 2
$ sudo pip install --upgrade $TF_BINARY_URL
```



<http://solutionsreview.com/cloud-platforms/microsoft-beats-aws-google-on-cloud-storage-benchmark-test/>



AWS GPU price in Oregon

GPU Instances - Current Generation

g2.2xlarge	8	26	15	60 SSD	\$0.65 per Hour
g2.8xlarge	32	104	60	2 x 120 SSD	\$2.6 per Hour

EC2 Console: Oregon

The screenshot shows the AWS EC2 console interface. At the top, there's a navigation bar with the AWS logo, a dropdown for 'AWS', a 'Services' dropdown, an 'Edit' dropdown, and user information for 'Sung Kim' in the 'Oregon' region. On the left, a sidebar menu lists 'EC2 Dashboard', 'Events', 'Tags', 'Reports', 'Limits', 'INSTANCES' (with 'Instances', 'Spot Requests', 'Reserved Instances', 'Scheduled Instances', 'Commands', 'Dedicated Hosts'), 'IMAGES' (with 'AMIs', 'Bundle Tasks'), and 'ELASTIC BLOCK STORE' (with 'Volumes', 'Snapshots'). The main content area has a title 'Resources' and a message: 'You are using the following Amazon EC2 resources in the US West (Oregon) region:'. It displays the following resource counts:

Resource Type	Count
Running Instances	0
Dedicated Hosts	0
Volumes	1
Key Pairs	2
Placement Groups	0
Elastic IPs	0
Snapshots	1
Load Balancers	0
Security Groups	13

Below this is a callout box with the text: 'Easily deploy Ruby, PHP, Java, .NET, Python, Node.js & Docker applications with [Elastic Beanstalk](#)'. The 'Create Instance' section contains a note: 'To start using Amazon EC2 you will want to launch a virtual server, known as an Amazon EC2 instance.' It features a prominent blue 'Launch Instance' button. A note below states: 'Note: Your instances will launch in the US West (Oregon) region'. The 'Service Health' and 'Scheduled Events' sections are also visible at the bottom.

Resources

You are using the following Amazon EC2 resources in the US West (Oregon) region:

Resource Type	Count
Running Instances	0
Dedicated Hosts	0
Volumes	1
Key Pairs	2
Placement Groups	0
Elastic IPs	0
Snapshots	1
Load Balancers	0
Security Groups	13

Easily deploy Ruby, PHP, Java, .NET, Python, Node.js & Docker applications with [Elastic Beanstalk](#).

Create Instance

To start using Amazon EC2 you will want to launch a virtual server, known as an Amazon EC2 instance.

[Launch Instance](#)

Note: Your instances will launch in the US West (Oregon) region

Service Health

Scheduled Events

Account Attributes

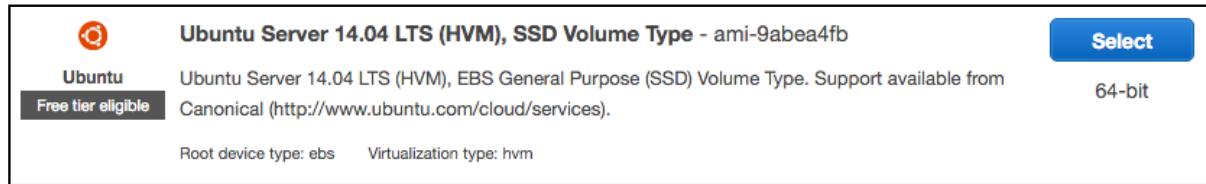
Supported Platforms
VPC
Default VPC
vpc-9b61f2fe
Resource ID length management

Additional Information

Getting Started Guide
Documentation
All EC2 Resources
Forums
Pricing
Contact Us

AWS Marketplace

Find free software trial products



1. Choose AMI 2. Choose Instance Type 3. Configure Instance 4. Add Storage 5. Tag Instance 6. Configure Security Group 7. Review

Step 2: Choose an Instance Type

Amazon EC2 provides a wide selection of instance types optimized to fit different use cases. Instances are virtual servers that can run applications. They have varying combinations of CPU, memory, storage, and networking capacity, and give you the flexibility to choose the appropriate mix of resources for your applications. [Learn more about instance types and how they can meet your computing needs.](#)

Filter by: GPU instances ▾ Current generation ▾ Show/Hide Columns

Currently selected: g2.2xlarge (26 ECUs, 8 vCPUs, 2.6 GHz, Intel Xeon E5-2670, 15 GiB memory, 1 x 60 GiB Storage Capacity)

	Family	Type	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS-Optimized Available	Network Performance
<input checked="" type="checkbox"/>	GPU instances	g2.2xlarge	8	15	1 x 60 (SSD)	Yes	High
<input type="checkbox"/>	GPU instances	g2.8xlarge	32	60	2 x 120 (SSD)	-	10 Gigabit

Step 4: Add Storage

Your instance will be launched with the following storage device settings. You can attach additional EBS volumes and instance store volumes to your instance, or edit the settings of the root volume. You can also attach additional EBS volumes after launching an instance, but not instance store volumes. [Learn more about storage options in Amazon EC2.](#)

Volume Type	Device	Snapshot	Size (GiB)	Volume Type	IOPS	Delete on Termination	Encrypted
Root	/dev/sda1	snap-306df873	12	General Purpose SSD	36 / 3000	<input checked="" type="checkbox"/>	Not Encrypted
Instance Store 0	/dev/sdb	N/A	N/A	N/A	N/A	<input type="checkbox"/>	Not Encrypted

Add New Volume

ubuntu,
GPU,
12G or more

key to access the server

Select an existing key pair or create a new key pair X

A key pair consists of a **public key** that AWS stores, and a **private key file** that you store. Together, they allow you to connect to your instance securely. For Windows AMIs, the private key file is required to obtain the password used to log into your instance. For Linux AMIs, the private key file allows you to securely SSH into your instance.

Note: The selected key pair will be added to the set of keys authorized for this instance. Learn more about [removing existing key pairs from a public AMI](#).

Choose an existing key pair

Select a key pair

hunkim-oregon

I acknowledge that I have access to the selected private key file (hunkim-oregon.pem), and that without this file, I won't be able to log into my instance.

[Cancel](#) [Launch Instances](#)

EC2: Create an instance

The screenshot shows the AWS EC2 Instances page. The left sidebar navigation includes:

- EC2 Dashboard
- Events
- Tags
- Reports
- Limits
- INSTANCES** (selected)
- Instances** (selected)
- Spot Requests
- Reserved Instances
- Scheduled Instances
- Commands
- Dedicated Hosts
- IMAGES**
- AMIs
- Bundle Tasks
- ELASTIC BLOCK STORE**
- Volumes
- Snapshots
- NETWORK & SECURITY**
- Security Groups
- Elastic IPs
- Placement Groups

The main content area displays a table of instances. The table has columns: Name, Instance ID, Instance Type, Availability Zone, Instance State, Status Checks, and Alarm Status. One instance is listed:

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status
i-4ad0c48d	g2.2xlarge	us-west-2b	shutting-down	None		

Below the table, a detailed view for the instance i-4ad0c48d is shown. It includes tabs for Description, Status Checks, Monitoring, and Tags. The Description tab displays the following information:

Instance ID	i-4ad0c48d	Public DNS	-
Instance state	shutting-down	Public IP	-
Instance type	g2.2xlarge	Elastic IP	-
Private DNS	-	Availability zone	us-west-2b

It's ready to ssh!

The screenshot shows the AWS EC2 Instances page. At the top, there is a search bar with the placeholder "Filter by tags and attributes or search by keyword". Below the search bar is a header row with columns: Name, Instance ID, Instance Type, Availability Zone, Instance State, Status Checks, Alarm Status, and Public DNS. A single instance is listed in the table:

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS
	i-068c65db	g2.2xlarge	us-west-2c	running	Initializing	None	ec2-54-186-153-9.us-w...

Below the table, the instance details are displayed:

Instance: i-068c65db Public DNS: ec2-54-186-153-9.us-west-2.compute.amazonaws.com

Buttons for Description, Status Checks, Monitoring, and Tags are shown below the instance details.

Instance ID	i-068c65db	Public DNS	ec2-54-186-153-9.us-west-2.compute.amazonaws.com
Instance state	running	Public IP	54.186.153.9
Instance type	g2.2xlarge	Elastic IP	-
Private DNS	ip-172-31-15-155.us-west-2.compute.internal	Availability zone	us-west-2c

Requires CUDA and CuDNN

```
# Ubuntu/Linux 64-bit, GPU enabled, Python 2.7
# Requires CUDA toolkit 7.5 and CuDNN v4. For other versions, see "Install from
sources" below.
$ export TF_BINARY_URL=https://storage.googleapis.com/tensorflow/linux/gpu/
tensorflow-0.9.0rc0-cp27-none-linux_x86_64.whl

# Python 2
$ sudo pip install --upgrade $TF_BINARY_URL
```

CUDA

The screenshot shows a blog post on a website. At the top, there's a header with the title "Seven Story Rabbit Hole" and a subtitle "Sometimes awesome things happen in deep rabbit holes. Or not." To the right of the subtitle is a cartoon illustration of a brown rabbit with long ears. Below the header, the navigation bar includes links for "Blog" and "Archives", and a search bar with a magnifying glass icon. The main content area features a date "NOV 22ND, 2015" and a large, bold title: "CUDA 7.5 on AWS GPU Instance Running Ubuntu 14.04". Underneath the title, there's a section titled "Launch stock Ubuntu AMI" with a bulleted list of instructions: "Launch ami-d05e75b8", "Choose a GPU instance type: g2.2xlarge or g2.8xlarge", and "Increase the size of the storage (this depends on what else you plan to install, I'd suggest at least 20 GB)". Below this, another section is titled "SSH in" with a dark blue code block containing the command: "1 \$ ssh ubuntu@<instance ip>". To the right of the main content, there's a sidebar titled "About Me" which describes the author as a software engineer at Couchbase working on Couchbase Mobile. It also mentions their work on ElasticThought and Caffe Deep Learning toolkit, and provides a Twitter handle "@tleydn". The sidebar also lists recent posts: "Setting Up a Self-hosted drone.io CI Server", "Adding Vendoring to a Go Project", "Configure Emacs as a Go Editor From Scratch Part 3", and "Octopress Under Docker".

<http://tleyden.github.io/blog/2015/11/22/cuda-7-dot-5-on-aws-gpu-instance-running-ubuntu-14-dot-04/>

cuDNN

Optional: cuDNN

One can apply for the developer program here

<https://developer.nvidia.com/cudnn>. When approved, download cuDNN for Linux (either v4 RC or v3 is fine), upload the cuDNN package from the local computer to the instance, and install cuDNN:

```
tar -zxf cudnn-7.0-linux-x64-v4.0-rc.tgz #or cudnn-7.0-linux-x64-v3.0
cd cuda
sudo cp lib64/* /usr/local/cuda/lib64/
sudo cp include/cudnn.h /usr/local/cuda/include/
```

<https://no2147483647.wordpress.com/2016/01/16/setup-amazon-aws-gpu-instance-with-mxnet/>

16 commands

- 1 wget http://developer.download.nvidia.com/.../cuda-repo-ubuntu1404...
- 2 sudo dpkg -i cuda-repo-ubuntu1404_7.5-18_amd64.deb
- 3 sudo apt-get update
- 4 sudo apt-get upgrade -y
- 5 sudo apt-get install -y opencl-headers build-essential protobuf-compiler libprotobuf-dev libboost-all-dev libleveldb-dev hdf5-tools libhdf5-serial-dev libopencv-core-dev libopencv-highgui-dev libsnavy-dev libsnavy1 libatlas-base-dev cmake libstdc++6-4.8-dbg libgoogle-glog0 libgoogle-glog-dev libgflags-dev liblmdb-dev git python-pip gfortran
- 6 sudo apt-get clean
- 7 sudo apt-get install -y linux-image-extra-`uname -r` linux-headers-`uname -r` linux-image-`uname -r`
- 8 sudo apt-get install -y cuda
- 9 nvidia-smi
- 10 sudo apt-get install python-pip python-dev
- 11 sudo pip install --upgrade https://storage.googleapis.com/.../tensorflow-0.8.0rc0-cp27-n...
- 12 git clone https://github.com/nlintz/TensorFlow-Tutorials
- 13 cd TensorFlow-Tutorials/
- 14 vi ~/.profile # add PATH, LD PATH
- 15 source ~/.profile
- 16 python 06_autoencoder.py

Add Path

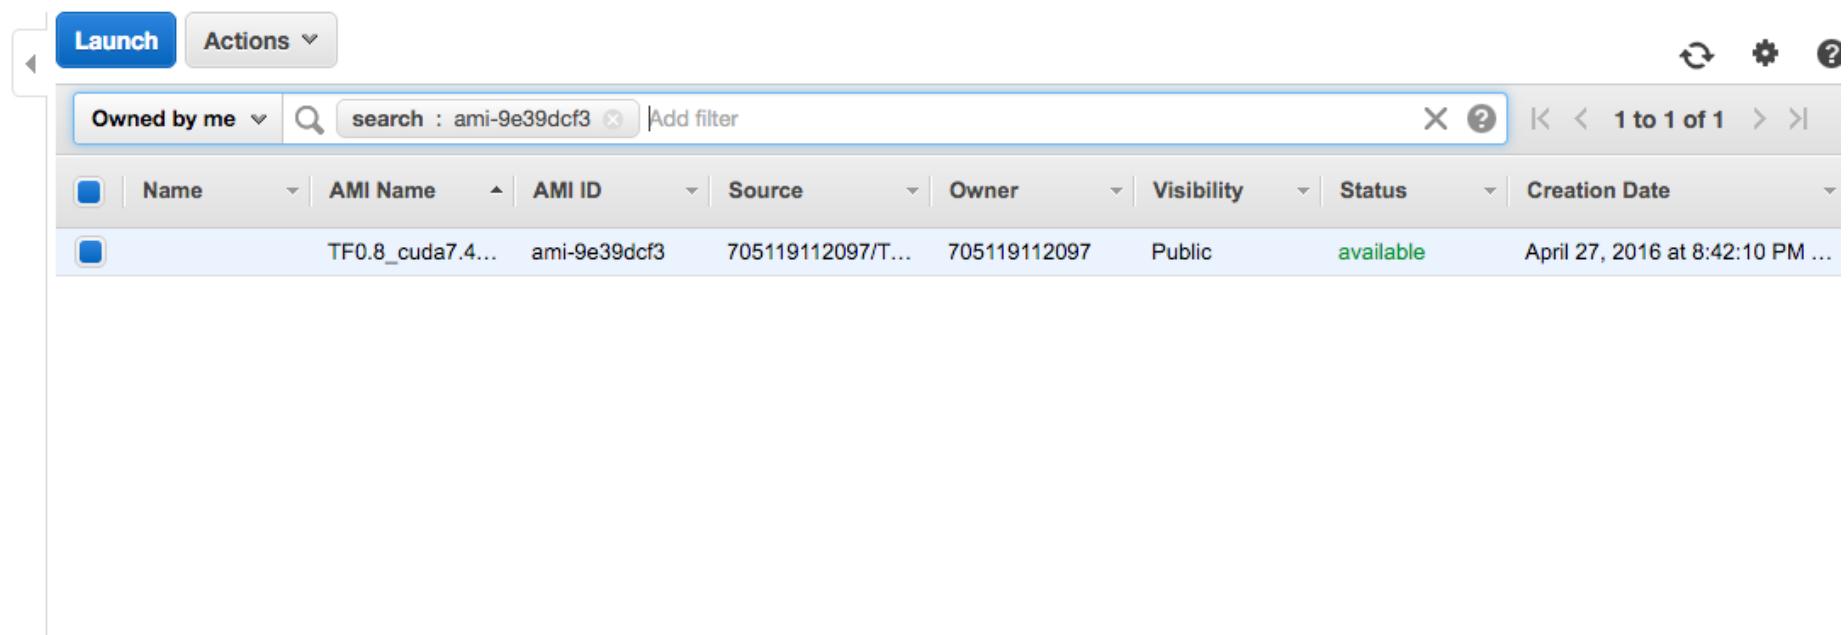
- `export PATH=/usr/local/cuda/bin:$PATH`
- `export LD_LIBRARY_PATH=/usr/local/cuda/lib64:$LD_LIBRARY_PATH`

Requires CUDA and CuDNN

```
# Ubuntu/Linux 64-bit, GPU enabled, Python 2.7
# Requires CUDA toolkit 7.5 and CuDNN v4. For other versions, see "Install from
sources" below.
$ export TF_BINARY_URL=https://storage.googleapis.com/tensorflow/linux/gpu/
tensorflow-0.9.0rc0-cp27-none-linux_x86_64.whl

# Python 2
$ sudo pip install --upgrade $TF_BINARY_URL
```

Reuse ami-9e39dcf3 (N.Virginia) ami-38f60658 (oregon)



The screenshot shows the AWS EC2 Dashboard with the 'AMIs' section selected. The main view displays a table of AMIs. A single row is visible, representing the AMI 'TF0.8_cuda7.4...' with ID 'ami-9e39dcf3'. The table includes columns for Name, AMI Name, AMI ID, Source, Owner, Visibility, Status, and Creation Date. The 'Status' column shows 'available' in green. The 'Creation Date' column shows 'April 27, 2016 at 8:42:10 PM ...'. The search bar at the top is set to 'Owned by me' and contains the search term 'ami-9e39dcf3'.

Name	AMI Name	AMI ID	Source	Owner	Visibility	Status	Creation Date
TF0.8_cuda7.4...	ami-9e39dcf3	705119112097/T...	705119112097	Public	available	April 27, 2016 at 8:42:10 PM ...	

Perhaps, it will be unavailable when later CUDA versions are out.

Creating TensorFlow device (/gpu:0)

```
05] successfully opened CUDA library libcublas.so locally
9] Couldn't open CUDA library libcudnn.so. LD_LIBRARY_PATH: /usr/local/cuda/lib64:
c:1562] Unable to load cuDNN DSO
05] successfully opened CUDA library libcufft.so locally
05] successfully opened CUDA library libcuda.so.1 locally
05] successfully opened CUDA library libcurand.so locally
-ubyte.gz')
cationWarning: converting an array with ndim > 0 to an index will result in an error in the
e]
.py:42: VisibleDeprecationWarning: converting an array with ndim > 0 to an index will result
, 1)
-ubyte.gz')
ubyte.gz')
ubyte.gz')
xecutor.cc:900] successful NUMA node read from SysFS had negative value (-1), but there must
t.cc:102] Found device 0 with properties:

7

t.cc:126] DMA: 0
t.cc:136] 0:   Y
ice.cc:755] Creating TensorFlow device (/gpu:0) -> (device: 0, name: GRID K520, pci bus id:
locator.cc:244] PoolAllocator: After 1704 get requests, put_count=1321 evicted_count=1000 ev
locator.cc:256] Raising pool_size_limit_ from 100 to 110
locator.cc:244] PoolAllocator: After 1704 get requests, put_count=1812 evicted_count=1000 ev
locator.cc:256] Raising pool_size_limit_ from 256 to 281
```

```
482/6750 (epoch 3), train_loss = 7.503, time/batch = 0.237
483/6750 (epoch 3), train_loss = 7.612, time/batch = 0.239
484/6750 (epoch 3), train_loss = 7.473, time/batch = 0.248
485/6750 (epoch 3), train_loss = 7.603, time/batch = 0.241
486/6750 (epoch 3), train_loss = 7.519, time/batch = 0.248
487/6750 (epoch 3), train_loss = 7.526, time/batch = 0.238
488/6750 (epoch 3), train_loss = 7.356, time/batch = 0.243
489/6750 (epoch 3), train_loss = 7.550, time/batch = 0.239
490/6750 (epoch 3), train_loss = 7.440, time/batch = 0.243
491/6750 (epoch 3), train_loss = 7.514, time/batch = 0.240
492/6750 (epoch 3), train_loss = 7.514, time/batch = 0.242
493/6750 (epoch 3), train_loss = 7.467, time/batch = 0.245
494/6750 (epoch 3), train_loss = 7.351, time/batch = 0.239
495/6750 (epoch 3), train_loss = 7.553, time/batch = 0.245
496/6750 (epoch 3), train_loss = 7.373, time/batch = 0.240
497/6750 (epoch 3), train_loss = 7.493, time/batch = 0.242
498/6750 (epoch 3), train_loss = 7.445, time/batch = 0.243
499/6750 (epoch 3), train_loss = 7.432, time/batch = 0.240
500/6750 (epoch 3), train_loss = 7.476, time/batch = 0.246
501/6750 (epoch 3), train_loss = 7.463, time/batch = 0.238
502/6750 (epoch 3), train_loss = 7.477, time/batch = 0.241
503/6750 (epoch 3), train_loss = 7.495, time/batch = 0.244
504/6750 (epoch 3), train_loss = 7.543, time/batch = 0.240
505/6750 (epoch 3), train_loss = 7.550, time/batch = 0.241
506/6750 (epoch 3), train_loss = 7.567, time/batch = 0.242
507/6750 (epoch 3), train_loss = 7.415, time/batch = 0.247
508/6750 (epoch 3), train_loss = 7.414, time/batch = 0.236
509/6750 (epoch 3), train_loss = 7.540, time/batch = 0.242
510/6750 (epoch 3), train_loss = 7.402, time/batch = 0.245
511/6750 (epoch 3), train_loss = 7.548, time/batch = 0.241
512/6750 (epoch 3), train_loss = 7.385, time/batch = 0.242
513/6750 (epoch 3), train_loss = 7.563, time/batch = 0.238
514/6750 (epoch 3), train_loss = 7.488, time/batch = 0.241
515/6750 (epoch 3), train_loss = 7.504, time/batch = 0.245
516/6750 (epoch 3), train_loss = 7.546, time/batch = 0.243
517/6750 (epoch 3), train_loss = 7.521, time/batch = 0.244
518/6750 (epoch 3), train_loss = 7.384, time/batch = 0.246
519/6750 (epoch 3), train_loss = 7.402, time/batch = 0.242
520/6750 (epoch 3), train_loss = 7.552, time/batch = 0.244
521/6750 (epoch 3), train_loss = 7.499, time/batch = 0.250
522/6750 (epoch 3), train_loss = 7.342, time/batch = 0.239
523/6750 (epoch 3), train_loss = 7.378, time/batch = 0.237
524/6750 (epoch 3), train_loss = 7.391, time/batch = 0.242
525/6750 (epoch 3), train_loss = 7.348, time/batch = 0.242
526/6750 (epoch 3), train_loss = 7.270, time/batch = 0.242
527/6750 (epoch 3), train_loss = 7.381, time/batch = 0.244
```

```
model saved to save/model.ckpt
1/6750 (epoch 0), train_loss = 11.127, time/batch = 4.697
2/6750 (epoch 0), train_loss = 11.176, time/batch = 4.430
3/6750 (epoch 0), train_loss = 10.946, time/batch = 4.604
4/6750 (epoch 0), train_loss = 10.616, time/batch = 4.483
5/6750 (epoch 0), train_loss = 10.207, time/batch = 5.354
6/6750 (epoch 0), train_loss = 9.723, time/batch = 5.674
7/6750 (epoch 0), train_loss = 9.449, time/batch = 5.972
8/6750 (epoch 0), train_loss = 9.089, time/batch = 5.785
9/6750 (epoch 0), train_loss = 8.861, time/batch = 5.518
10/6750 (epoch 0), train_loss = 8.797, time/batch = 5.856
11/6750 (epoch 0), train_loss = 8.688, time/batch = 5.795
12/6750 (epoch 0), train_loss = 8.814, time/batch = 5.803
13/6750 (epoch 0), train_loss = 8.728, time/batch = 5.595
14/6750 (epoch 0), train_loss = 8.938, time/batch = 5.586
15/6750 (epoch 0), train_loss = 8.657, time/batch = 5.529
16/6750 (epoch 0), train_loss = 8.736, time/batch = 5.536
17/6750 (epoch 0), train_loss = 8.643, time/batch = 5.690
18/6750 (epoch 0), train_loss = 8.985, time/batch = 5.828
19/6750 (epoch 0), train_loss = 8.792, time/batch = 5.504
20/6750 (epoch 0), train_loss = 8.882, time/batch = 5.555
21/6750 (epoch 0), train_loss = 8.750, time/batch = 5.437
22/6750 (epoch 0), train_loss = 8.872, time/batch = 5.597
```

□



OS X El Capitan

Version 10.11.3

MacBook Pro (Retina, 13-inch, Early 2015)

Processor 3.1 GHz Intel Core i7

Memory 16 GB 1867 MHz DDR3

AWS GPU price in Oregon

GPU Instances - Current Generation					
g2.2xlarge	8	26	15	60 SSD	\$0.65 per Hour
g2.8xlarge	32	104	60	2 x 120 SSD	\$2.6 per Hour

Amazon Elastic Compute Cloud running Linux/UNIX

\$0.650 per On Demand Linux g2.2xlarge Instance Hour	17 Hrs	\$11.05
\$2.6 per On Demand Linux g2.8xlarge Instance Hour	1 Hrs	\$2.60
Total:		\$13.65

$$2.6 * 24 * 30 = 1,872 \text{ USD}$$

Spot instances

The screenshot shows the AWS EC2 Dashboard with the 'Spot Requests' tab selected. The main area displays a message: "You do not have any Spot instance requests in this region." Below this, a call-to-action says "Click on the 'Request Spot Instances' button to request your first Spot instance." A large blue button labeled "Request Spot Instances" is prominently displayed. The top navigation bar includes links for AWS services like Lambda, Step Functions, and CloudWatch, along with user information for Sung Kim and the N. Virginia region.

AWS Services Edit Sung Kim N. Virginia

EC2 Dashboard Events Tags Reports Limits

INSTANCES Instances **Spot Requests** Reserved Instances

Request Spot Instances Cancel Pricing History

Filter by tags and attributes or search by keyword

You do not have any Spot instance requests in this region.

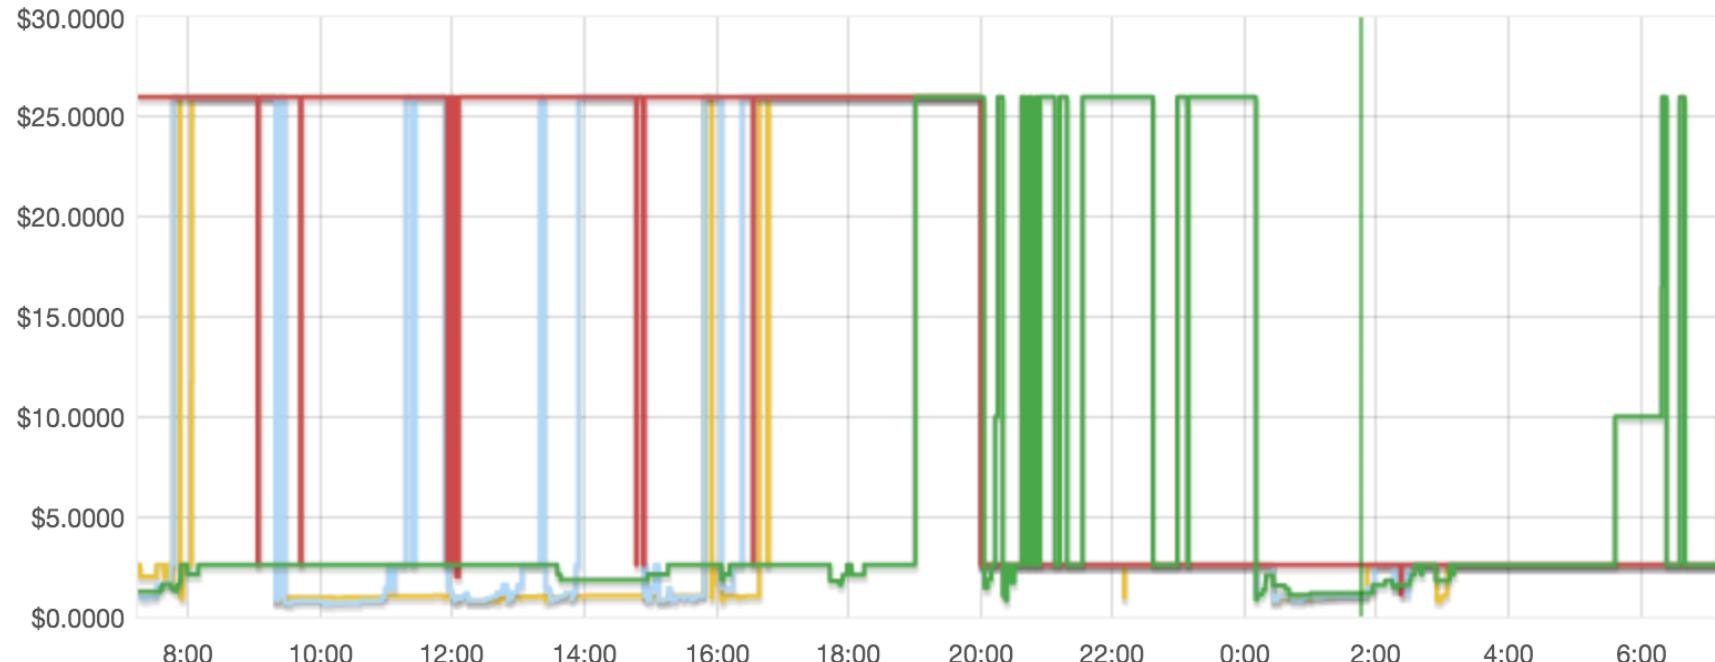
Click on the "Request Spot Instances" button to request your first Spot instance.

Request Spot Instances

Spot Instance Pricing History

X

Product : [Linux/UNIX](#) ▾ Instance type: [g2.8xlarge](#) ▾ Date range : [1 day](#) ▾ Availability zone: [All zones](#) ▾



Availability zone	Price
us-east-1a	\$1.0485
us-east-1c	\$1.0200
us-east-1d	\$2.6001

N. Virginia ▾



None fo

- EC2 Dashboard
- Events
- Tags
- Reports
- Limits
- INSTANCES
- Instances
- Spot Requests**
- Reserved Inst

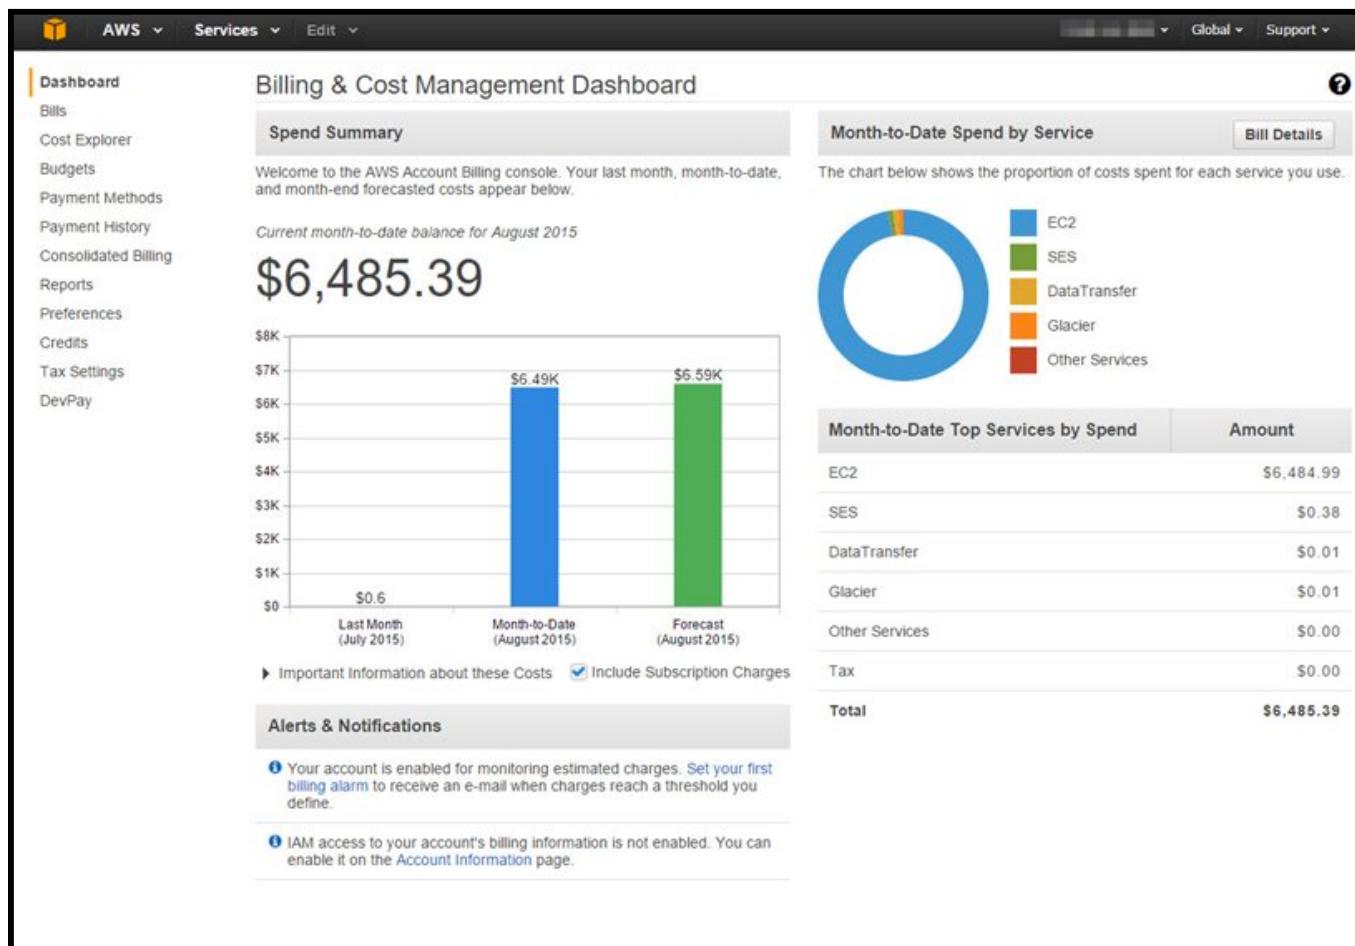
Price bidding

Step 3: Configure Instance Details

Configure the instance to suit your requirements. You can launch multiple instances from the same AMI, request Spot instances to take advantage of the lower pricing, assign an access management role to the instance, and more.

Number of instances i	<input type="text" value="1"/>	Launch into Auto Scaling Group i								
Purchasing option i	<input checked="" type="checkbox"/> Request Spot instances									
Current price i	<table><tbody><tr><td>us-east-1a</td><td>0.2428</td></tr><tr><td>us-east-1c</td><td>0.1442</td></tr><tr><td>us-east-1d</td><td>0.153</td></tr><tr><td>us-east-1e</td><td>6.500</td></tr></tbody></table>		us-east-1a	0.2428	us-east-1c	0.1442	us-east-1d	0.153	us-east-1e	6.500
us-east-1a	0.2428									
us-east-1c	0.1442									
us-east-1d	0.153									
us-east-1e	6.500									
Maximum price i	\$ <input type="text" value="(e.g. 0.045 = 4.5 cents/hour)"/>									
Launch group i	(Optional)									

bill, bill, bill!



Check, stop, and terminate

The screenshot shows the AWS EC2 Dashboard. On the left, there's a sidebar with links like EC2 Dashboard, Events, Tags, Reports, Limits, Instances (which is expanded to show Instances and Spot Requests), and EC2 Metrics. The main area is titled "Resources" and displays the following statistics for the US West (Oregon) region:

Category	Value
Running Instances	0
Dedicated Hosts	0
Volumes	0
Key Pairs	2
Placement Groups	0
Elastic IPs	0
Snapshots	1
Load Balancers	0
Security Groups	14

On the right, under "Account Attributes", it shows the following information:

- Supported Platforms: VPC
- Default VPC: vpc-9b61f2fe
- Resource ID length management

The screenshot shows the AWS EC2 Instances page. At the top, there are buttons for "Launch Instance", "Connect", and "Actions". A dropdown menu is open under "Actions" with the following options: Connect, Get Windows Password, Launch More Like This, Instance State, Instance Settings, Image, Networking, and CloudWatch Monitoring. The "Instance State" option is highlighted. Below the dropdown, a message says "You do not have any running instances in this region." and "Want to learn more about using EC2? Check out the [Getting Started Guide](#)." To the left, there's a note "Select an instance above".

Cloud Watch

The screenshot shows the AWS EC2 Instances page. On the left, there's a sidebar with links like EC2 Dashboard, Events, Tags, Reports, Limits, Instances (which is selected), Spot Requests, Reserved Instances, and Images. The main area shows a table of instances. One instance, 'kbill', is selected. An 'Actions' dropdown menu is open over this instance, showing options: Connect (Get Windows Password), Launch More Like This (Instance State, Instance Settings, Image, Networking), CloudWatch Monitoring (Enable Detailed Monitoring, Disable Detailed Monitoring, Add/Edit Alarms), and a 'Filter by tags and attributes or search' bar.

This is a modal dialog titled 'Alarm Details for i-7858ecdf (kbill)'. It contains the following text: 'Below are your CloudWatch alarms for the selected resources. Click on an alarm to edit it or click 'view' to see additional options and details in Amazon CloudWatch.' Below this, a table header is shown with columns for State, Name, and More Options. The body of the dialog states 'There are no alarms for this instance'. At the bottom, there's a 'Create Alarm' button and a 'Close' button.

<http://docs.aws.amazon.com/AmazonCloudWatch/latest/DeveloperGuide/UsingAlarmActions.html#AddingStopActions>

Stop when CPU utilization ≤ 0.3

Create Alarm X

You can use CloudWatch alarms to be notified automatically whenever metric data reaches a level you define.

To edit an alarm, first choose whom to notify and then define when the notification should be sent.

Send a notification to: No SNS topics found... ▼

Take the action: Recover this instance (i)
 Stop this instance (i)
 Terminate this instance (i)
 Reboot this instance (i)

Whenever: Average ▼ of CPU Utilization ▼

Is: \leq Percent

For at least: consecutive period(s) of ▼

CPU Utilization Percent



4/30 18:00 20:00 22:00

Name of alarm: awsec2-i-7858ecdf-High-CPU-Utilization

[Cancel](#) **Create Alarm**

Shutdown after training

```
$ screen
```

```
$ sudo -i
```

```
# python train.py; shutdown -h now
```

Deep learning for Everyone

Season 2 coming soon!

Sung Kim <hunkim+ml@gmail.com>
<http://hunkim.github.io/ml/>

